

Data Science in Python

Basic data visualization

Matplotlib in a nutshell

Matplotlib is the basic Python plotting library

- Generate publication-quality visualizations
- Export results in pdf, svg, png, jpg
- Mostly intended for *static* 2D graphs
- Rather low-level: you build everything from scratch

Today's modules

```
import matplotlib.pyplot as plt # standard matplotlib import
import numpy as np
```

Custom style options

```
plt.style.use('seaborn-v0_8-whitegrid')
plt.rcParams['figure.figsize'] = [8, 2.5]
plt.rcParams['font.size'] = 14
```

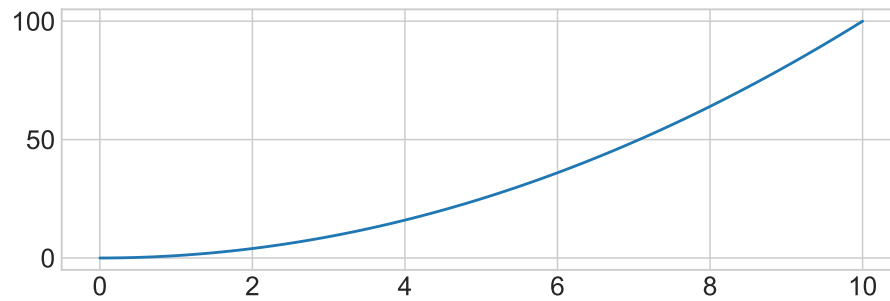
Getting started

Generate some data: $y = x^2$

```
x = np.linspace(0, 10, 100) # 100 evenly spaced points in [0, 10]
y = x**2 # x squared - a parabola!
```

Let's make a simple **line plot** of y against x

```
fig = plt.figure() # Figure object
plt.plot(x, y) # horizontal, vertical
```

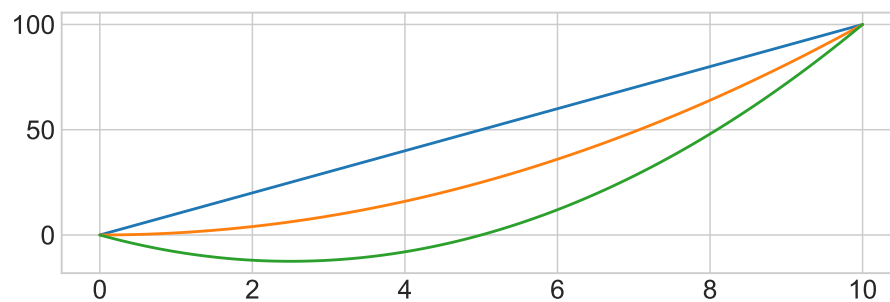


Multiple lines

We can plot multiple curves in the same figure

```
x = np.linspace(0, 10, 100) # 100 points evenly spaced in [0, 10]
y1 = 10*x # linear
y2 = x**2 # quadratic
y3 = 2*x**2 - 10*x # another quadratic
```

```
fig = plt.figure()
plt.plot(x, y1)
plt.plot(x, y2)
plt.plot(x, y3)
```



Don't like the default style? Look at the several options available!

Matplotlib options

There are several options to customize:

- Figure size
- Line color, width, and style
- Marker size, color, and type
- Axes labels

- Legend
- Axes range
-

Check out <https://matplotlib.org> and the [gallery](#) to discover all the options.

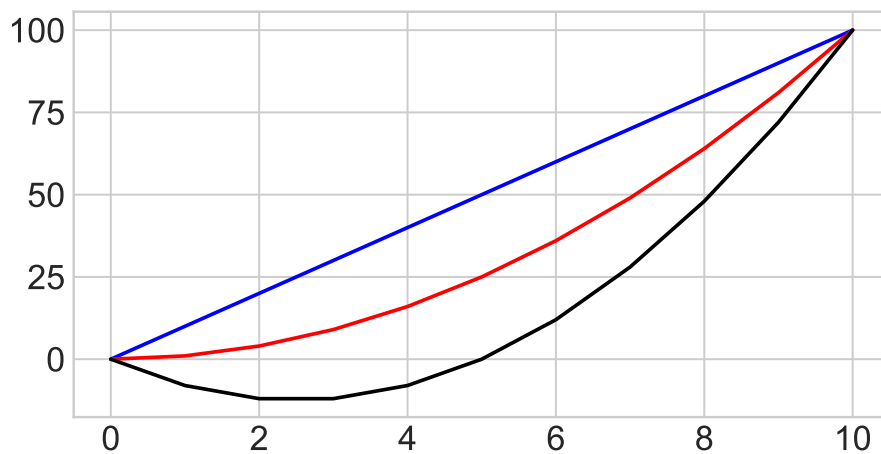
Style options: line color

A *style option string* is often used to change the appearance of a plotted object, such as the color:

- `b`: $\Rightarrow$ Blue, `r`: $\Rightarrow$ Red, `k`: $\Rightarrow$ black
- `g`: Green, `m`: Magenta, `c`: Cyan, `y`: Yellow

```
x = np.arange(11) # 0, 1, 2, ... 10
y1 = 10*x; y2 = x**2; y3 = 2*x**2 - 10*x

plt.figure(figsize=(6, 3)) # specify figure size
plt.plot(x, y1, 'b') # plot in Blue color
plt.plot(x, y2, 'r') # plot in Red color
plt.plot(x, y3, 'k') # plot in black color
plt.grid(True) # show a grid
```

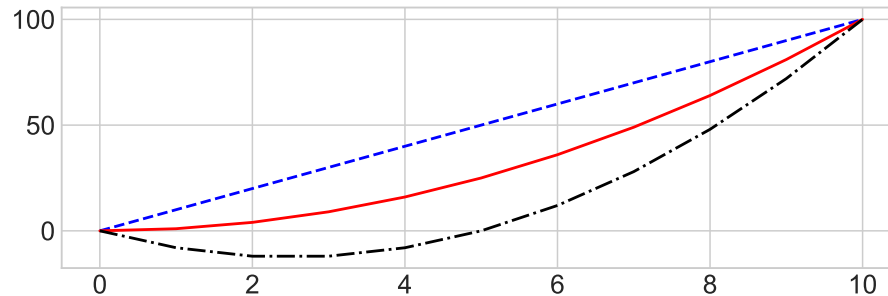


Style options: line style

We can also specify the line type with style options:

- `'b--'` $\Rightarrow$ Blue color, dashed line
- `'r-'` $\Rightarrow$ Red color, continuous line
- `'k-.'` $\Rightarrow$ black color, dashed-dotted line

```
plt.figure()
plt.plot(x, y1, 'b--') # plot in Blue color, dashed line
plt.plot(x, y2, 'r-') # plot in Red color, continuous line
plt.plot(x, y3, 'k-.') # plot in black color, dashed-dotted line
plt.grid(True) # show/hide grid
```



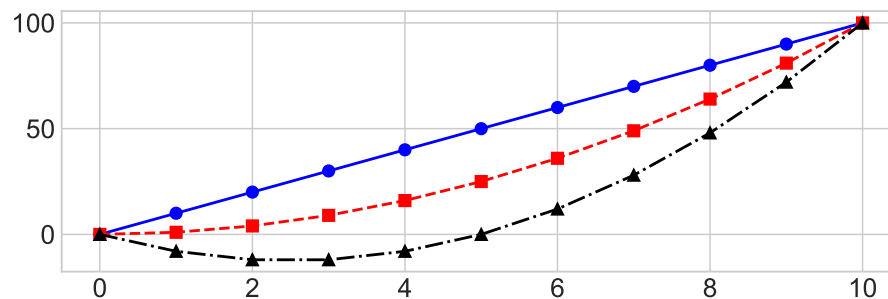
We combine line color and style in a *style string* option. Old convention taken from MATLAB.

Style options: marker

We can also specify the marker type:

- 'bo-' ⇒ Blue color, marker circle, continuous line
- 'rs--' ⇒ Red color, marker square, dashed line
- 'k^-. ' ⇒ black color, marker triangle, dashed-dotted line

```
plt.figure()
plt.plot(x, y1, 'bo-') # plot in Blue color, continuous line, with circles as markers.
plt.plot(x, y2, 'rs--') # plot in Red color, dashed line, square as markers.
plt.plot(x, y3, 'k^-.'); # plot in black color, dashed-dotted line
```

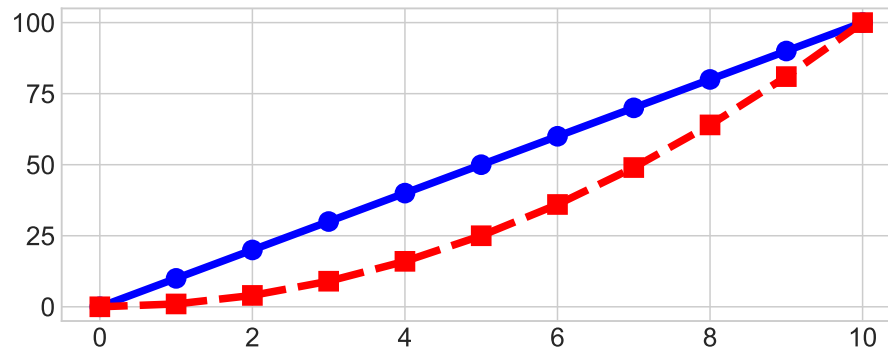


Other style options

Certain style options can only be specified as extra arguments. Here we change:

- Line width
- Marker size
- ...

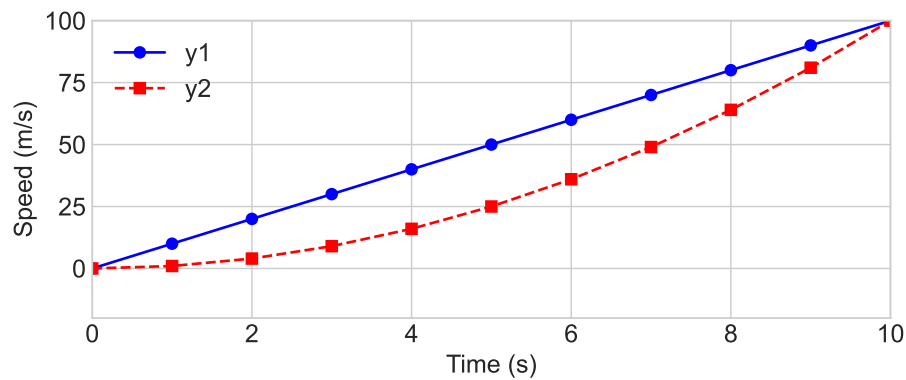
```
plt.figure(figsize=(8, 3))
plt.plot(x, y1, 'bo-', linewidth=4, markersize=10) # Blue color, continuous line, with circles
plt.plot(x, y2, 'rs--', linewidth=4, markersize=10); # Red color, dashed line, square as marker
```



Again, check out the matplotlib [documentation](#)!

Axes limits, grid, and legend

```
plt.figure(figsize=(8, 3.0))
# blue, continuous line, circles as markers, name y1.
plt.plot(x, y1, 'bo-', label='y1')
plt.plot(x, y2, 'rs--', label='y2')
plt.xlabel('Time (s)') # label for the x axis
plt.ylabel('Speed (m/s)') # label for the y axis
plt.xlim([0, 10]) # range of the x axis
plt.ylim([-20, 100]) # range of the y axis
plt.legend() # show a legend
plt.grid(True) # show a grid
```

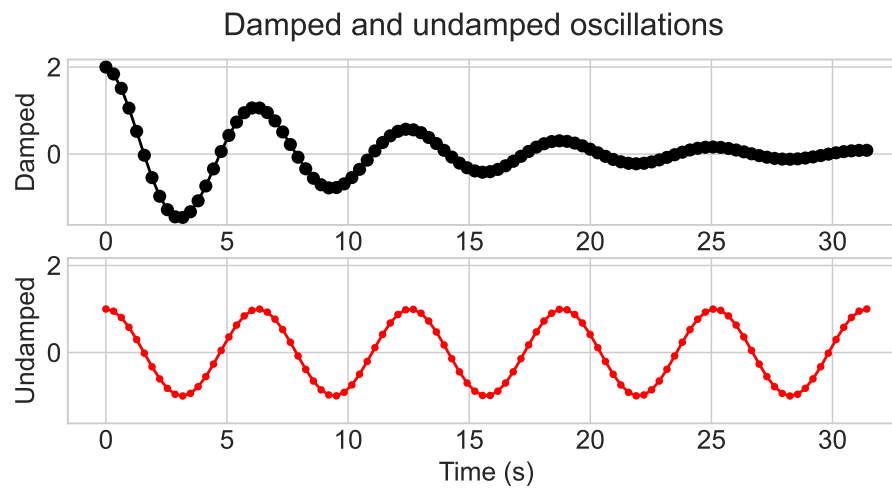


- We can specify:
 - *limits* for the horizontal and vertical axes
 - *labels* for the horizontal and vertical axes
 - *labels* for the different plotted objects, to be shown in a *legend*
- We can turn the *gridline* on and off
- ...

Subplots

Show several *panels* horizontally/vertically aligned. Example from the [matplotlib gallery](#):

```
# 2 panels aligned vertically, sharing the same y axis
fig, ax = plt.subplots(2, sharey=True, figsize=(8, 3.5))
fig.suptitle("Damped and undamped oscillations")
ax[0].plot(x, y1, 'ko-')
ax[0].set_ylabel('Damped')
ax[1].plot(x, y2, 'r.-')
ax[1].set_ylabel('Undamped')
ax[1].set_xlabel('Time (s)');
```

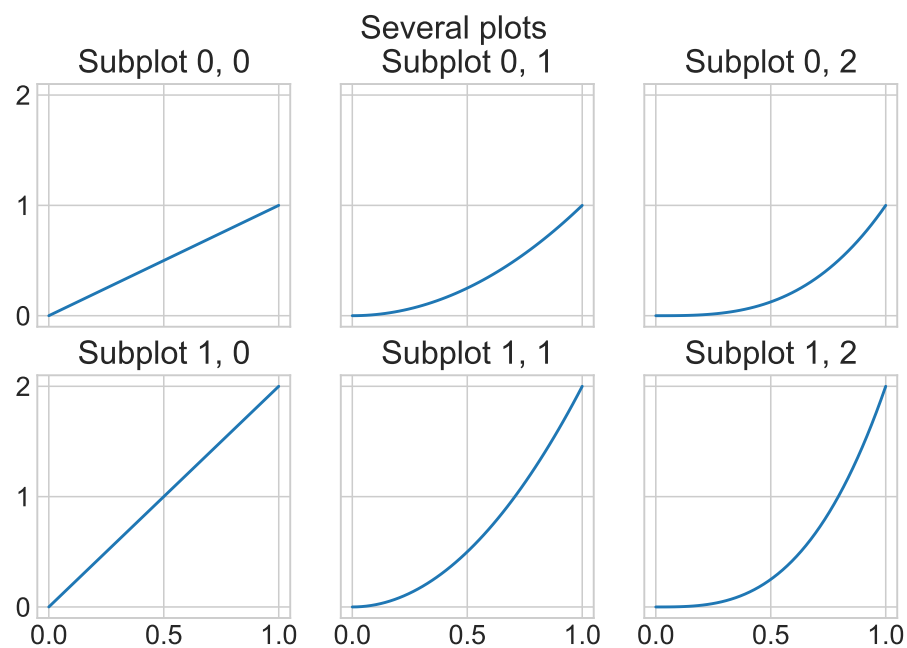


Note: the two plots have the same scale on the y axis (option `sharey=True`)

Subplots

An example with rows and columns of panels

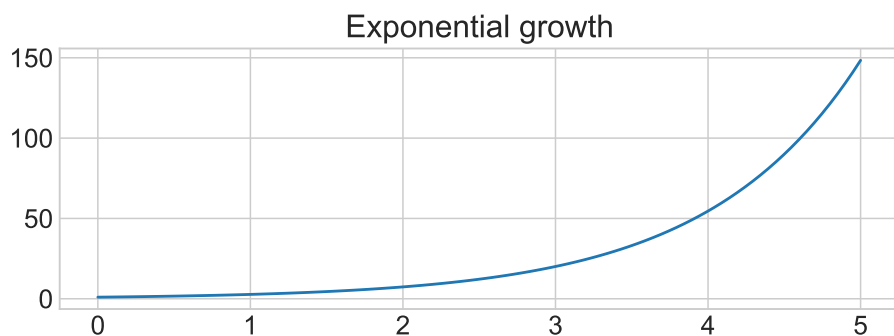
```
n_row = 2; n_col = 3; x = np.linspace(0.0, 1.0)
fig, ax = plt.subplots(n_row, n_col, sharex=True, sharey=True, figsize=(8, 5)) # 2x3 subplot
fig.suptitle("Several plots")
for row_idx in range(n_row):
    for col_idx in range(n_col):
        ax[row_idx][col_idx].plot(x, (row_idx+1) * x**(col_idx+1))
        ax[row_idx][col_idx].set_title(f"Subplot {row_idx}, {col_idx}")
```



Save figure to file

```
x = np.linspace(0, 5, 100) # 100 points in [0, 5]
y = np.exp(x)

fig = plt.figure()
plt.title("Exponential growth")
plt.plot(x, y)
plt.savefig("exponential.png")
```



Matplotlib figures can be exported to common formats such as .png, .jpeg, .pdf,

.svg.

Exercises (until Subplots included)

Commonly used plots

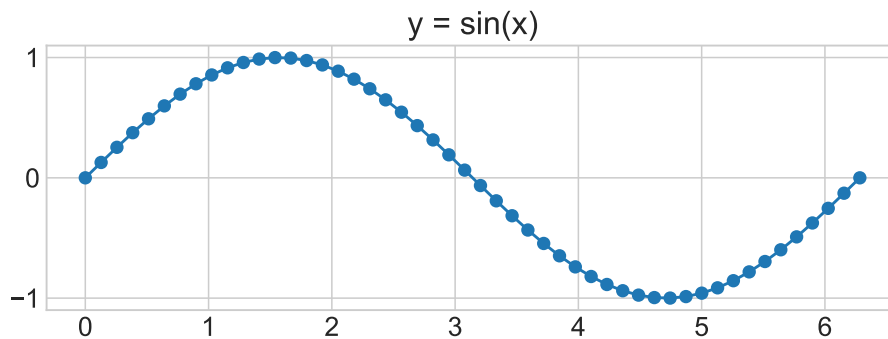
- Line plot
- Scatter plot
- Bar plot
- Histogram
- Boxplot

Line plot

A line plot is used to represent a continuous function $y = f(x)$. It uses *interpolation* to produce a continuous smooth curve.

```
x = np.linspace(0, 2*np.pi, 50)
y = np.sin(x)

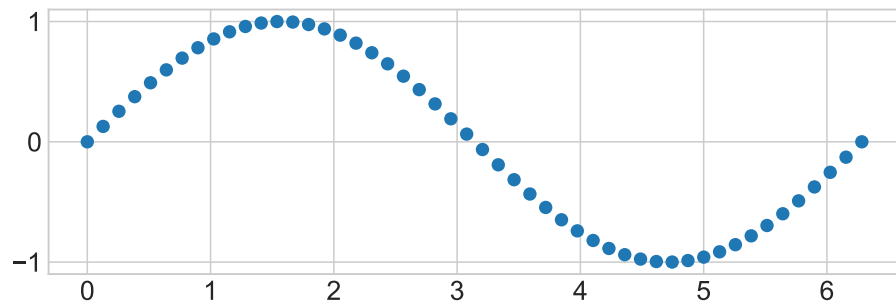
plt.figure()
plt.plot(x, y, 'o-')
plt.title("y = sin(x)");
```



Scatter plot

A scatter plot is used to display data points in 2D. Like a lineplot, without interpolation.

```
x = np.linspace(0, 2*np.pi, 50)
y = np.sin(x)
plt.plot(x, y, 'o'); # use circles as markers
```



Can be implemented using the same plot function, using a *marker* style option such as 'o'.

Scatter plot cont'd

The specialized `scatter` functions allows setting *data-dependent* marker color & size.

```
n_points = 20
x = np.random.randn(n_points)
y = np.random.randn(n_points)
z = np.random.randn(n_points)
w = np.random.rand(n_points)

plt.figure(figsize=(8, 4))
plt.scatter(x, y, c=z, s=w*2000, alpha=0.3, cmap='viridis')
plt.colorbar(); # show color scale
```



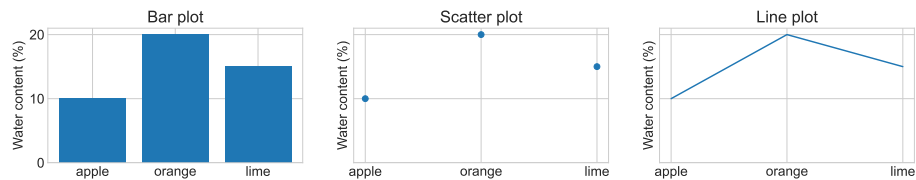
This enables exploration of 4 dimensions at once!

Bar plot

More intuitive than a scatter/line plot when the quantity on the x axis is categorical (i.e. not numeric)

```
names = ['apple', 'orange', 'lime']; values = [10, 20, 15]

fig, ax = plt.subplots(1, 3, figsize=(16, 2.5), sharey=True)
ax[0].set_title("Bar plot")
ax[0].bar(names, values)
ax[0].set_ylabel("Water content (%)")
ax[1].set_title("Scatter plot")
ax[1].scatter(names, values)
ax[1].set_ylabel("Water content (%)")
ax[2].set_title("Line plot")
ax[2].plot(names, values);
ax[2].set_ylabel("Water content (%)");
```



The line plot is particularly misleading in this example: how to interpolate apples and oranges?

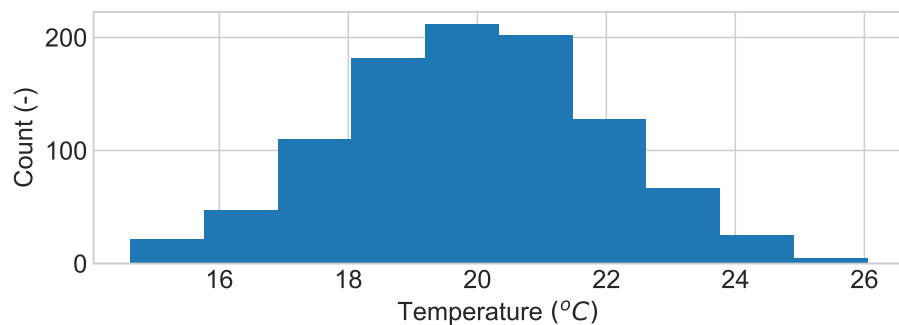
Histogram

A histogram visually describes the *empirical distribution* of one numerical variable.

- The numerical range of a variable is divided into *bins*.
- For each bin, the number of times the variable is within the bin range is counted
- The bins may be either user-defined or automatically constructed

```
temperatures = 20 + 2 * np.random.randn(1000) # temperatures for different days

plt.figure()
plt.hist(temperatures)
plt.xlabel('Temperature  $(^{\circ}\text{C})$ ')
plt.ylabel('Count (-)');
```

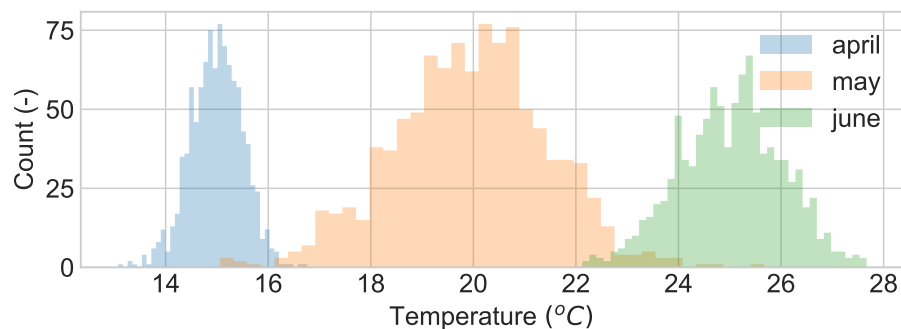


Histogram cont'd

A fancier example with multiple histograms and a legend

```
# fake temperature data for three months
temp_april = np.random.normal(15, 0.5, 1000) # mean, std, n_samples
temp_may = np.random.normal(20, 1.5, 1000)
temp_june = np.random.normal(25, 1, 1000)

kwargs = dict(histtype='stepfilled', alpha=0.3, bins=40)
plt.hist(temp_april, **kwargs, label='april')
plt.hist(temp_may, **kwargs, label='may')
plt.hist(temp_june, **kwargs, label='june')
plt.xlabel('Temperature $(^{\circ}\text{C})$')
plt.ylabel('Count (-)');
plt.legend();
```



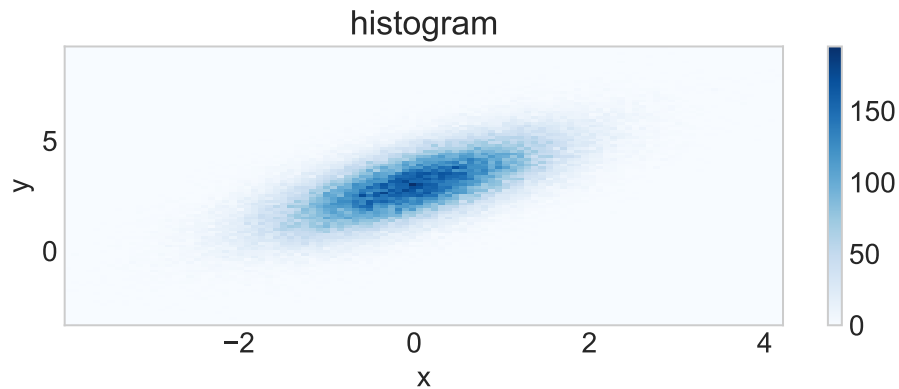
Overlap handled here with *transparency* (**alpha** channel option)

Histogram in 2D

Histograms may be extended to 2-dimensional point clouds. Use the `hist2d` function

```
mean = [0, 3]; cov = [[1, 1], [1, 2]]
x, y = np.random.multivariate_normal(mean, cov, 100000).T

plt.hist2d(x, y, bins=100, cmap='Blues')
plt.colorbar()
plt.title("histogram")
plt.xlabel("x")
plt.ylabel("y");
```



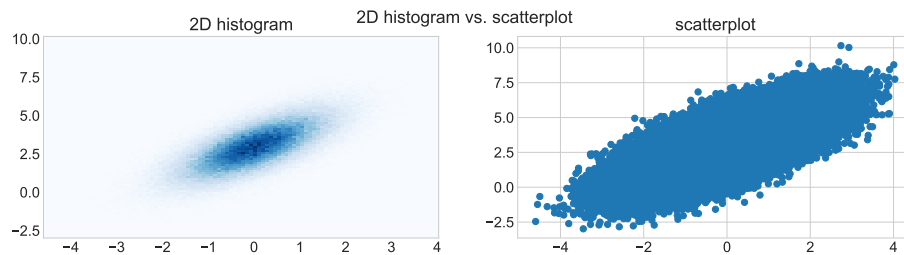
The histogram count is represented using a *color map*. Here, dark blue is high density.

Histogram in 2D cont'd

Histograms may be extended to 2-dimensional point clouds. Use the `hist2d` function

```
mean = [0, 3]; cov = [[1, 1], [1, 2]];
x, y = np.random.multivariate_normal(mean, cov, 200000).T

fig, ax = plt.subplots(1, 2, figsize=(15, 3.5))
fig.suptitle("2D histogram vs. scatterplot")
ax[0].hist2d(x, y, bins=100, cmap='Blues')
ax[0].set_title("2D histogram")
ax[1].scatter(x, y)
ax[1].set_title("scatterplot");
```

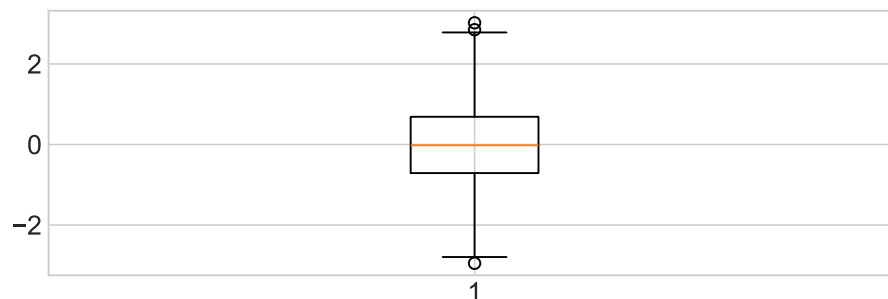


The 2D histogram (left) may be more informative than a 2D scatterplot (right).

Boxplot

A box plot summarizes the salient aspects of an empirical distribution.

```
data = np.random.randn(1000) # random data point
plt.boxplot(data);
```



- The orange line is the *median* of the distribution
 - 50% of the points are inside the box
 - Most of the points are expected to be within the interval defined by the *wiskers*
 - Potential *outliers* are shown as circles
-

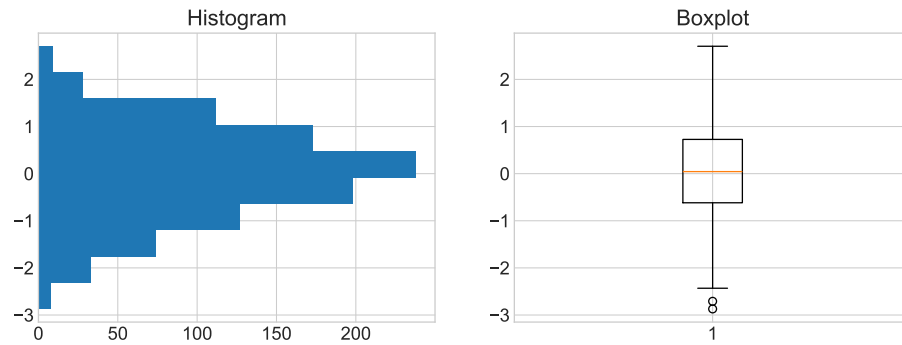
Boxplot cont'd

It is less informative than a histogram, but perhaps easier to read. Less is (sometimes) more!

```
data = np.random.randn(1000)

fig, ax = plt.subplots(1, 2, figsize=(12, 4))
ax[0].hist(data, orientation="horizontal")
```

```
ax[0].set_title("Histogram")
ax[1].boxplot(data)
ax[1].set_title("Boxplot");
```



Boxplot cont'd

We often look at several boxplots simultaneously.

```
temp_april = np.random.normal(15, 0.5, 1000) # mean, std, n_samples
temp_may = np.random.normal(20, 1.5, 1000)
temp_june = np.random.normal(25, 1, 1000)

plt.boxplot([temp_april, temp_may, temp_june], tick_labels = ['april', 'may', 'june'])
plt.ylabel("Temperature ( $^{\circ}\text{C}$ )");
```

<>:6: SyntaxWarning:


invalid escape sequence '\m'

<>:6: SyntaxWarning:

invalid escape sequence '\m'

/var/folders/w7/h67z1rn941j1p0m51k08d9xc0000gp/T/ipykernel_7198/1002276405.py:6: SyntaxWarning:

invalid escape sequence '\m'



- Typical visualization for datasets with a **mix of continuous and categorical** variables
- Analyses on mixed continuous/categorical variables **required in the project!**
 - Choose a dataset that support such analyses!

Final Presentation:

How to find a dataset? Start thinking! Warning: **plagiarism will not be tolerated !**

Where to find a dataset?

- Your hobbies and interests (games, sports, travels...)
 - Competitions often have (un)official datasets/APIs (e.g. Formula 1's [Fast-F1](#))
- Data collected on your own, e.g. with your smartphone

See next slides...

-
- Multi-Domain Portals & Aggregators
 - Kaggle Datasets: Thousands of community-contributed datasets across every domain. Most are tabular (CSV). <https://www.kaggle.com/datasets>
 - UCI Machine Learning Repository: Classic benchmark datasets for classification, regression, and clustering. Most datasets are tabular with accompanying documentation. <https://archive.ics.uci.edu/datasets>
 - Google Dataset Search: Meta-search engine indexing datasets from thousands of repositories worldwide. Useful for discovering niche or domain-specific data. <https://datasetsearch.research.google.com/>
 - Hugging Face Datasets: Large collection of NLP, vision, and audio datasets. Accessible via the *datasets* Python library. <https://huggingface.co/datasets>
 - EU Open Data Portal: Official open data portal of the European Union. Contains datasets from EU institutions on economy, environ-

- ment, transport, agriculture, and public health. Many datasets are tabular (CSV/Excel). <https://data.europa.eu/en>
- opendata.swiss: Official open data portal of the Swiss Confederation. Over 10000 datasets from federal, cantonal, and communal authorities covering transport, environment, health, economy, education, and public administration. Most data is tabular (CSV/JSON). <https://opendata.swiss/en/>
 - Data.gov (USA): U.S. government open data portal with over 250000 datasets covering health, education, climate, finance, agriculture, and more. <https://data.gov/>
 - Awesome Public Datasets (GitHub): A curated list of high-quality public datasets organised by topic. Useful starting point for discovering data across many domains. <https://github.com/awesome-data/awesome-public-datasets>

- Swiss Datasets

- Swiss Federal Statistical Office (BFS/OFS): Comprehensive statistics on the Swiss population, economy, education, health, environment, and society. <https://www.bfs.admin.ch/bfs/en/home.html>
- MeteoSwiss Open Data: Historical and real-time meteorological data for Switzerland: temperature, precipitation, sunshine hours, wind, and more from over 260 stations. <https://www.meteoswiss.admin.ch/services-and-publications/applications/ext/download-data-without-registration.html>
- SBB Open Data (Swiss Federal Railways): Real-time and historical data on Swiss public transport: train punctuality, passenger counts, station information, and timetables. <https://opentransport-data.swiss/en/>
- Swiss National Forest Inventory (LFI): Data on Swiss forests: tree species, volumes, health, and biodiversity from periodic national surveys. <https://www.lfi.ch/resultate/>
- Swiss Open Government Data – Canton of Zürich: Open data specifically from the Canton of Zürich: demographics, buildings, traffic, education, health statistics, and more. Tabular formats (CSV). <https://www.zh.ch/de/politik-staat/opendata.html>
- FOEN – Swiss Federal Office for the Environment: Environmental monitoring data: air quality (NABEL network), water levels, noise pollution, waste statistics, biodiversity indicators for Switzerland. <https://www.bafu.admin.ch/bafu/en/home/state/data.html>
- Swiss Personenverkehr (Public Transport Statistics): Detailed statistics on Swiss public transport usage, including modal split, passenger-kilometres, and regional comparisons. <https://www.bfs.admin.ch/bfs/en/home/statistics/mobility-transport.html>

- Weather, Climate & Environment
 - NOAA Climate Data Online: Global historical weather data: daily/monthly summaries from weather stations worldwide, including temperature, precipitation, wind, and extreme events. <https://www.ncei.noaa.gov/cdo-web/>
 - NASA Earthdata: Satellite remote sensing data: land surface temperature, vegetation indices (NDVI), sea surface temperature, ice sheet observations, and atmospheric data. <https://earthdata.nasa.gov/>
 - Global Historical Climatology Network (GHCN): Daily and monthly climate summaries from over 100000 stations worldwide. One of the most widely used datasets in climate science. <https://www.ncei.noaa.gov/products/land-based-station/global-historical-climatology-network-daily>
 - Berkeley Earth Surface Temperature: Land and ocean temperature records with statistical analysis. Tabular data in text/CSV format with global coverage from the 1750s onward. <https://berkeleyearth.org/data/>
 - Global Biodiversity Information Facility (GBIF): Over 2 billion species occurrence records with geolocation data. Covers plants, animals, fungi, and microorganisms worldwide. <https://www.gbif.org/>
 - European Flood Awareness System (EFAS): Flood forecasting and historical discharge data for European river basins, provided through the Copernicus Emergency Management Service. <https://www.efas.eu/>
-
- Finance & Economics
 - Yahoo Finance: Historical stock prices, volume, dividends, and splits for global equities, ETFs, indices, and currencies. Access via the Python *yfinance* library. <https://pypi.org/project/yfinance/>
 - FRED – Federal Reserve Economic Data: Over 800000 economic time series: GDP, unemployment, interest rates, inflation, housing, and monetary aggregates. <https://fred.stlouisfed.org/>
 - World Bank Open Data: Development indicators for all countries: GDP per capita, poverty rates, health expenditure, education, trade, and more. API and CSV/Excel download. <https://data.world-bank.org/>
 - IMF Data: International Monetary Fund datasets: World Economic Outlook, Balance of Payments, Government Finance Statistics, and exchange rates. <https://data.imf.org/>
 - Quandl / Nasdaq Data Link: Financial, economic, and alternative data. The free tier includes many datasets. <https://data.nasdaq.com/>
 - Swiss National Bank (SNB) Data Portal: Swiss monetary and financial statistics: interest rates, exchange rates, balance of payments, banking statistics, and securities data. <https://data.snb.ch/en>
-

Where to find a dataset?

- Health & Medical
 - WHO Global Health Observatory: Global health statistics: life expectancy, disease burden, immunisation, maternal and child health, communicable diseases, and health systems across all countries. <https://www.who.int/data/gho>
 - PhysioNet: Repository of physiological signal datasets: ECG, EEG, blood pressure waveforms, sleep studies, and clinical databases. <https://physionet.org/>
 - Our World in Data – Health: Curated datasets on global health topics: COVID-19, vaccinations, disease burden, mental health, substance use, nutrition, and mortality. CSV/GitHub download. <https://github.com/owid/owid-datasets>
 - Global Burden of Disease (GBD): Comprehensive data on mortality, morbidity, and risk factors for 369 diseases across 204 countries from the Institute for Health Metrics and Evaluation. <https://vizhub.healthdata.org/gbd-results/>
 - Swiss Health Observatory (Obsan): Swiss health system data: health status, healthcare utilisation, costs, chronic diseases, and mental health indicators for Switzerland. <https://www.obsan.admin.ch/en>

- Space & Astronomy
 - NASA Exoplanet Archive: Confirmed exoplanets and candidate detections with orbital parameters, stellar properties, and discovery methods. <https://exoplanetarchive.ipac.caltech.edu/>
 - ESA Gaia Archive: Positions, distances, and motions of nearly 2 billion stars in the Milky Way from ESA's Gaia mission. <https://gea.esac.esa.int/archive/>
 - NASA Open Data Portal: Datasets from across NASA missions: planetary science, heliophysics, astrophysics, Earth science, and aeronautics. <https://data.nasa.gov/>
 - Minor Planet Center (MPC): Orbital elements, observations, and designations for all known asteroids and comets. <https://www.minorplanetcenter.net/data>
 - Open Exoplanet Catalogue: Community-maintained XML/CSV catalogue of all known exoplanets with detailed parameters. https://github.com/OpenExoplanetCatalogue/open_exoplanet_catalogue

- Sports
 - FBref – Football (Soccer) Statistics: Comprehensive match and player statistics for major football leagues worldwide. Covers goals, assists, xG, passing, and defensive actions. <https://fbref.com/>

- Statsbomb Open Data: Free event-level football data for select competitions (FA Women’s Super League, La Liga, FIFA World Cup, etc.). JSON format with detailed match events. <https://github.com/statsbomb/open-data>
- Lahman Baseball Database: Complete batting, pitching, and fielding statistics for Major League Baseball from 1871 to present. CSV format. <https://www.seanlahman.com/baseball-archive/statistics/>
- NBA Stats API: Official NBA statistics: player and team stats, play-by-play, shot charts, and advanced metrics. Accessible via the *nba api* Python package. https://github.com/swar/nba_api
- Transfermarkt: Player market values, transfer histories, and performance data for football leagues worldwide. <https://www.transfermarkt.com/>
- FIFA World Cup Data (Kaggle): Historical data on all FIFA World Cup matches from 1930 onward: scores, teams, venues, attendance, and referee information. <https://www.kaggle.com/datasets/abecklas/fifa-world-cup>
- Swiss Football League Data: Match results, standings, and statistics for the Swiss Super League, available through the Swiss Football Association and football data aggregators. <https://www.football.ch/SFL/en.aspx>